



Distributed computing with Spark and Python

A. Zonca - San Diego Supercomputer Center

What is Spark?

- A distributed computing framework

Connect to a spark cluster

Running on 4 instances of SDSC Cloud

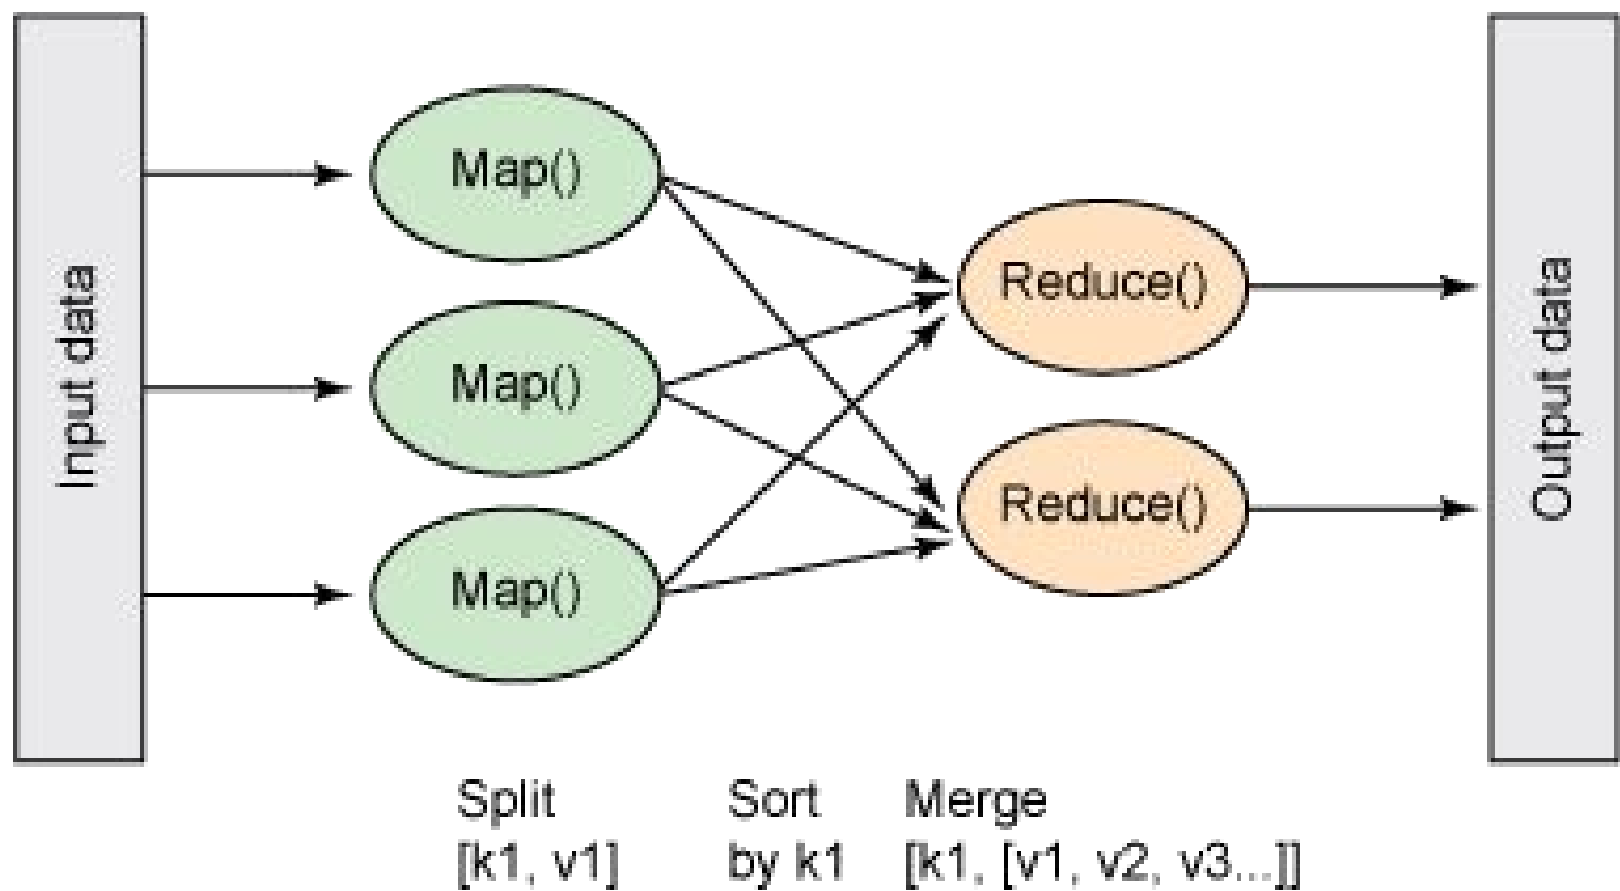
<http://bit.ly/sparknotebook>

password: acluster

ALWAYS COPY WITH YOUR NAME

House price example

see `house_price.ipynb`



map-reduce diagram

What's the point?

- write simple mapper and reducer
- framework scales to thousands of machines

Problem 1: Storage

- Big data
- Commodity hardware (Cloud)

Solution: Distributed File System

- redundant
- fault tolerant

Problem 2: Computation

- Slow to move data across network
- Computations fail

Solution: Hadoop Mapreduce / Spark

- Execute computation where data are located
- Rerun failed jobs

Problem 3: Communication

- Most of the times, need to summarize data to get a result
- Reduction phase in MapReduce
- Need data transfer across network

Solution: highly optimized Shuffle (All-to-All)

Spark and Hadoop

- Works within the Hadoop ecosystem
- Extends MapReduce
- Initially developed at UC Berkeley
- Now within the Apache Foundation
- ~400 and more developers


Key features of Spark

- **Resiliency:** tolerant to node failures
- **Speed:** supports in-memory caching
- **Ease of use:**
 - Python/Scala interfaces
 - interactive shells
 - many distributed primitives




	Hadoop MR Record	Spark Record	Spark 1 PB
Data Size	102.5 TB	100 TB	1000 TB
Elapsed Time	72 mins	23 mins	234 mins
# Nodes	2100	206	190
# Cores	50400 physical	6592 virtualized	6080 virtualized
Cluster disk throughput	3150 GB/s (est.)	618 GB/s	570 GB/s
Sort Benchmark Daytona Rules	Yes	Yes	No
Network	dedicated data center, 10Gbps	virtualized (EC2) 10Gbps network	virtualized (EC2) 10Gbps network
Sort rate	1.42 TB/min	4.27 TB/min	4.27 TB/min
Sort rate/node	0.67 GB/min	20.7 GB/min	22.5 GB/min

Spark 100TB benchmark

Worker Node






Worker Nodes






Cluster Manager
YARN/Standalone
Provision/Restart Workers

Worker Nodes

**Exec
JVM**

 Python
 Python
 Python

**Exec
JVM**

 Python
 Python
 Python

**Exec
JVM**

 Python
 Python
 Python

Worker Nodes

Driver Program

Spark
Context

Spark
Context

Cluster
Manager

Exec
JVM


Python
Python
Python

Exec
JVM

Python
Python
Python

Exec
JVM

Python
Python
Python



Spark Local

Driver Program

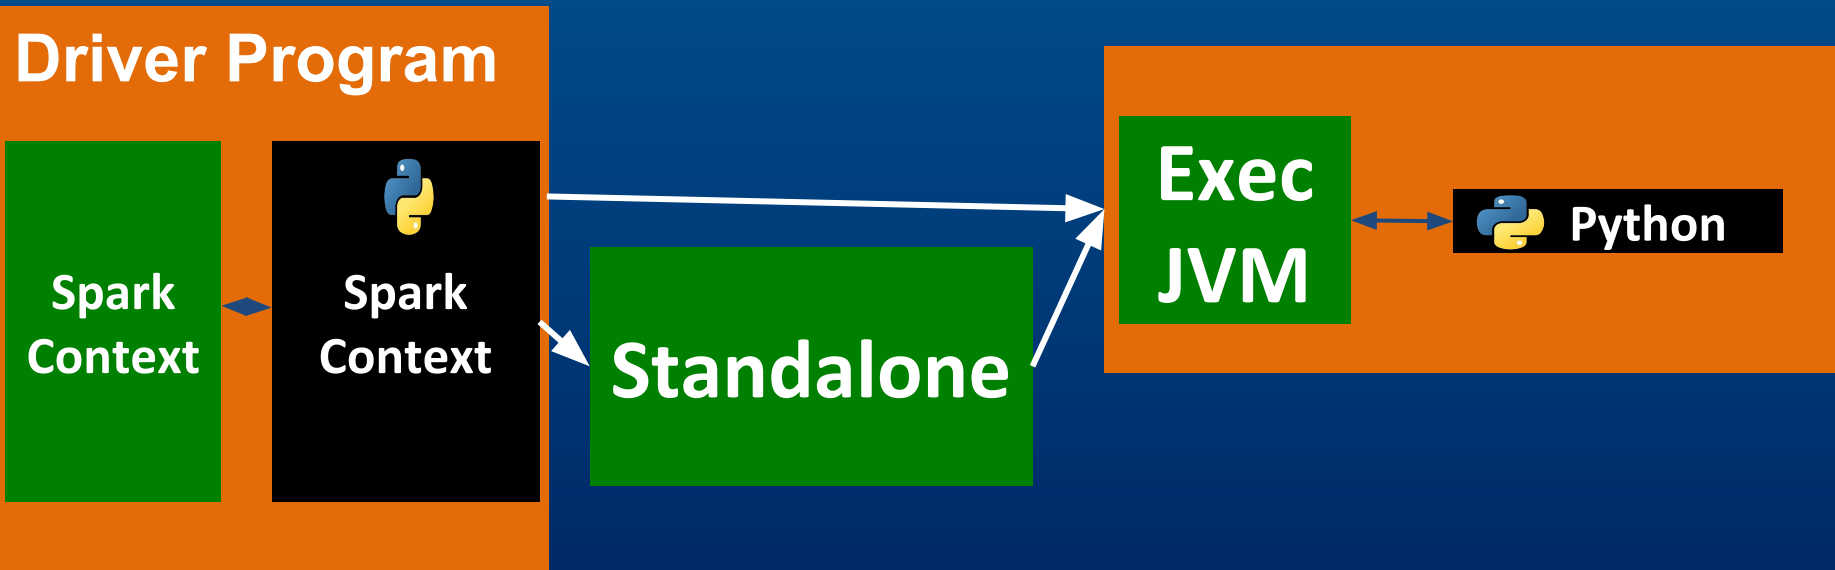
Spark
Context


Spark
Context

Standalone

Exec
JVM

 Python



on Comet

Master node

Driver Program

Spark
Context

Spark
Context

Standalone
CM

Computing Nodes

Exec
JVM


Python
Python
Python

Exec
JVM

Python
Python
Python

Exec
JVM


Python
Python
Python



on Amazon EMR

EC2 nodes

Master node



SDSC Cloud

Master node

Driver Program

Spark
Context

Spark
Context

Standalo
ne

Computing nodes

Exec
JVM



Python



Python



Python

Exec
JVM



Python



Python



Python

Exec
JVM



Python



Python



Python

House price with HDFS


see `house_price_hdfs.ipynb`

Resilient Distributed Dataset

- Resilient: fault tolerant, lineage is saved, lost partitions can be recovered
- Distributed: partitions are automatically distributed across nodes
- Created from: HDFS, S3, HBase, Local file, Local hierarchy of folders

map

map : apply function to each element of RDD




Other transformations

- `filter(func)` - keep only elements where func is true
- `sample(withReplacement, fraction, seed)` - get a random data fraction
- `coalesce(numPartitions)` - merge partitions to reduce them to numPartitions

filter

filter : keep only elements where func is true



coalesce

```
sc.parallelize(range(10), 4).glom().collect()
```


```
Out[]: [[0, 1], [2, 3], [4, 5], [6, 7, 8, 9]]
```

```
sc.parallelize(range(10), 4).coalesce(2).glom().collect()
```

```
Out[]: [[0, 1, 2, 3], [4, 5, 6, 7, 8, 9]]
```

coalesce


coalesce : reduce the number of partitions



Wide transformations

reduceByKey(func)

(K, V) pairs $\Rightarrow (K, \text{reduce } V \text{ with func})$



Narrow

vs

Wide

map

reduceByKey(sum)

A, 1

A, 3

A, 2

Wide transformations

- `groupByKey` : (K, V) pairs => (K, iterable of all V)
- `reduceByKey(func)` : (K, V) pairs => (K, result of reduction by func on all V)
- `repartition(numPartitions)` : similar to coalesce, shuffles all data to increase or decrease number of partitions to numPartitions

Shuffle


Shuffle

- . Global redistribution of data
- . High impact on performance

Shuffle

**requests
data over the
network**

**writes to
disk**



Know shuffle, avoid it


- . Which operations cause it?
- . Is it necessary?

Really need groupByKey?


`groupByKey`: (K, V) pairs => (K, iterable of all V)

if you plan to call `reduce` later in the pipeline,
use `reduceByKey` instead.

groupByKey + reduce



reduceByKey



Extract data from RDD


- `collect()` - copy all elements to the driver
- `take(n)` - copy first n elements
- `saveAsTextFile(filename)` - save to file
- `reduce(func)` - aggregate elements with func (takes 2 elements, returns 1)

Cached RDD

- Generally recommended after data cleaning
- Reusing cached data: 10x speedup
- Great for iterative algorithms
- If RDD too large, will only be partially cached in memory

Directed Acyclic Graph Scheduler

Directed Acyclic Graphs



Directed Acyclic Graphs

Track dependencies!
(also known as lineage or
provenance)

DAG in Spark

- nodes are RDDs
- arrows are Transformations



Narrow

vs

Wide

map

groupByKey





Narrow

vs

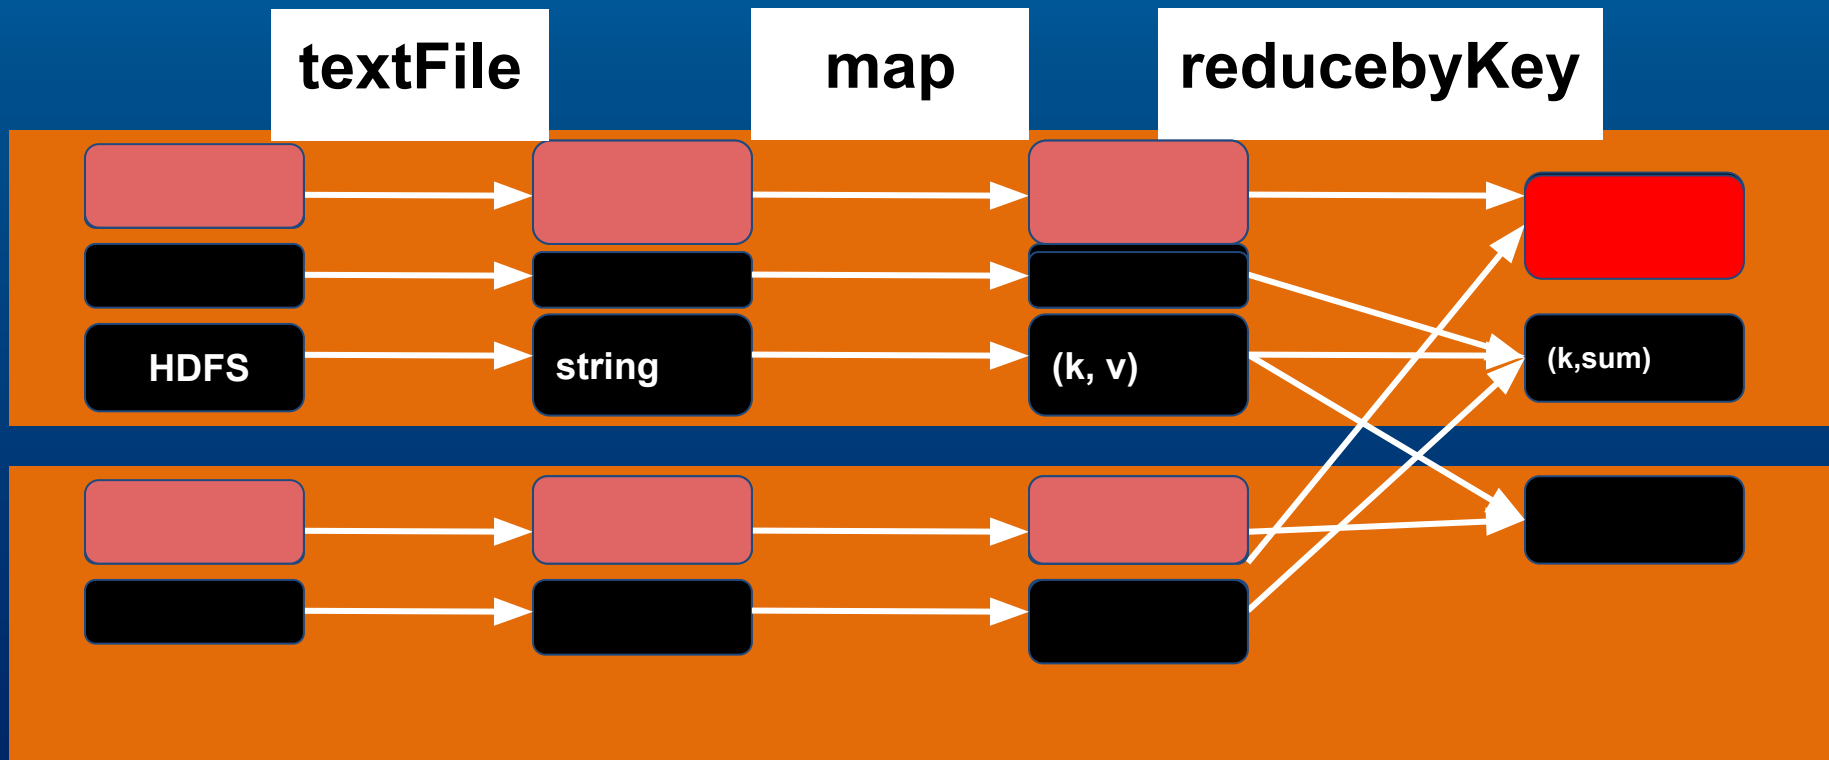
Wide

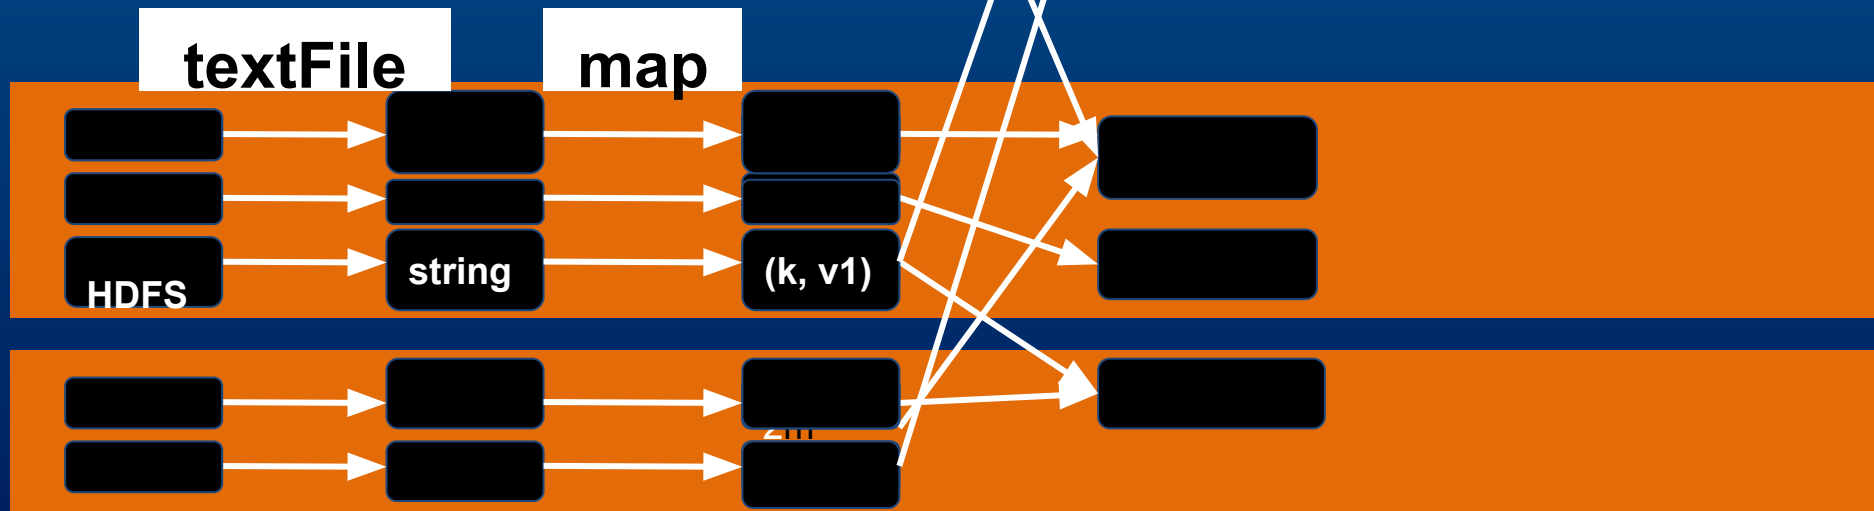

map

groupByKey



Spark DAG of transformations





Thank you

zonca@sdsc.edu